

Architecture Document

Introduction

Retrieval-Augmented Generation (RAG) is an emerging tech that combines information retrieval with large language models (LLMs) to produce context-aware, grounded answers. Unlike pure LLM-based approaches, RAG systems leverage external knowledge bases, ensuring outputs are factually accurate, domain-specific, and less prone to hallucination. This project implements a modular, Naive RAG system as a baseline to support enhancements, experimentation, and evaluation in subsequent phases.

The system is designed around three pillars: embedding, retrieval, and generation. It employs all-MiniLM-L6-v2 and all-mpnet-base-v2 embeddings from Sentence-Transformers, stores them in a vector database (Milvus Lite), and retrieves relevant chunks for query-conditioned generation. This document describes the architecture, design decisions, trade-offs, and extensibility considerations.

System Architecture

1. High-Level Workflow

Document Ingestion: Raw text is preprocessed, segmented into smaller chunks, and normalized.

Embedding Generation: Each chunk is embedded into a fixed-size dense vector representation using all-MiniLM-L6-v2.

Vector Indexing: Embeddings are inserted into Milvus Lite for similarity search.

Query Encoding: User queries are embedded into the same vector space.

Retrieval: Top-k most similar document chunks are retrieved via cosine similarity.

Response Generation: Retrieved evidence is injected into a prompt template, and an LLM generates the final answer.

2. System Components

Embedding Layer

Library: HuggingFace sentence-transformers.

Model: all-MiniLM-L6-v2 and all-mpnet-base-v2

Embedding sizes tested: 384 (default) and 784 (experimental).

Vector Database

Milvus Lite - scalable, distributed, production-oriented.

Retriever

Query embeddings compared with stored vectors using cosine similarity.

Configurable top-k retrieval (baseline = 1).


Generator

Takes retrieved document(s) and reformulates them into an answer.

Supports various prompting strategies (Instruction, Chain-of-Thought, Persona).

Design Decisions and Trade-offs

1. Embedding Model
all-MiniLM-L6-v2: 384 Embedding Size
All-mpnet-base-v2: 784 Embedding Size
2. Vector Store Choice:
Milvus Lite: Production-friendly, distributed retrieval, persistent storage.
Options: ChromaDB, FAISS
3. Retrieval Strategy
Baseline: Top k = 1
Experimentation: Top k = 3, 5



Documents --> Preprocess & Chunk --> Embedding Generation --> Vector DB

Query --> Query Encoding -----> Retriever -----> Generator --> Final Answer